

# CS 439: Intro to Data Science

## Course Project Guidelines

Created by Naina Chaturvedi

### 1 General Instruction

You should work on a project that is related to the theme of the course: Intro to Data Science. The scope of data science is broadly defined and can include any aspect of data collection, storage, integration, ML models, algorithms etc. I encourage you to pick a problem you are excited about and will be flexible if the project is relevant to topics and research papers in lectures.

### 2 Group Formation and Submit the Project Topic

You can form a project group (of max 3 people) and decide/submit the project topic by Nov 13<sup>th</sup>.

### 3 Project Guidelines

#### 3.1 Components

Your project should cover both Data Science and ML components.

**Project** - Formulate your own project question related to Intro to Data Science and attempt to provide a solution, which can be partial/complete with preliminary results. Your project can be computational, theoretical, experimental, or empirical. A project can be done individually (not recommended) or by a group of max 3 students (recommended, but no more than 3).

#### 3.2 Extra Notes

- **Start Early:** Start thinking about your project early and spend enough time developing it.

### 4 Deliverables

#### 4.1 1. Project Proposal

**Due:** 23:59pm, Nov 13<sup>th</sup>, 2025 (encouraged to submit early)

**Purpose:** The written proposal should define your project/research question and explain what you are planning to do.

**Length:** 2 pages (no more than 3), typed, single-space.

**Content:**

- **Define Project:** What problem are you solving? What strategic aspects are involved? How does your project relate to the lectures/papers we discussed?
- **Novelty and Importance:** Why is your project important? Why are you excited about it? What are some existing issues in current data science practices? Are there any prior related works? Provide a brief summary.
- **Plan:** Be specific and succinct.
  - What kind of data will you use (if any)? How will you get it? Will you create or simulate it?
  - What models/techniques/algorithms do you plan to use or develop?
  - What would be your implementation steps? How will you evaluate your method? How will you test and measure success?

## 4.2 2. Final Report

**Due:** 23:59pm, Dec 9<sup>th</sup>, 2025

**Length:** Normally, a well-explained project would take 6-8 pages, typed, single-space.

**Content:**

- **Project Definition:**
  - What problem are you solving? What strategic aspects are involved? How does your project relate to the lectures/papers we discussed?
- **Novelty and Importance:**
  - Why is your project important? Why are you excited about it? What are some existing issues in current data science practices? Are there any prior related works? Provide a brief summary.
  - Depending on your individual case, the above two aspects can be an extended or revised version of what you have written in your proposal.
- **Progress and Contribution:**
  - What kind of data did you use (if any)? How did you get it?
  - What models/techniques/algorithms did you use or develop?
  - What experiments did you design?
  - What are the key findings or results from your project? Did they verify or refute your original hypothesis? How did you evaluate your method?
  - Discuss the advantages and limitations of your approach.
- **Changes After Proposal:**
  - If your final report differs from your proposed project, discuss the differences, why you made certain changes, and the bottlenecks that prevented you from proceeding with the proposed project.

**Note to all:** You may use tools to help with your writing but do not use generated contents directly. Please cite any tools, web sources, papers, and textbooks you consult/use. You are responsible for the content of your writing, including its originality and correctness. Plagiarism is not allowed.

## 5 Project Suggestions/Examples

Refer to projects on Cubits (Links to be shared later).

## 6 Example Projects and Their Structure

**Note:** These are sample projects and their structure. You can define your own structure for your project.

### 6.1 1. Customer Churn Prediction for a Telecom Company

**Objective:** Predict customer churn using historical customer data and machine learning techniques.

**Steps and Technologies:**

#### A. Data Collection:

- Extract data from company's CRM system via API
- Collect customer reviews through web scraping
- Set up real-time data collection from service logs
- Implement data quality checks at collection points

#### B. Data Storage Strategy:

- Store structured data with proper organization
- Implement storage solutions for unstructured data (customer reviews, call logs)
- Create data archival procedures for historical data
- Organize data files with appropriate naming conventions

#### C. Data Integration:

- Create data pipelines for combining multiple data sources
- Implement data synchronization procedures
- Set up real-time data integration monitoring

- Use workflow management tools for automation

#### **D. Data Cleaning:**

- Handle missing values using Pandas
- Remove duplicates and standardize formats
- Create data quality reports
- Implement automated data cleansing procedures

#### **E. Data Transformation:**

- Create feature engineering pipelines
- Normalize and scale features
- Generate derived metrics for churn prediction
- Create aggregated views for analysis

#### **F. Exploratory Data Analysis:**

- Visualize customer behavior patterns using Matplotlib
- Create distribution analysis using Seaborn
- Generate correlation studies
- Build interactive dashboards

#### **G. Model Building:**

- Develop machine learning models (logistic regression, random forests)
- Implement cross-validation procedures
- Create model validation pipelines
- Set up model versioning

#### **H. Evaluation and Deployment:**

- Create model performance metrics
- Evaluate model accuracy and effectiveness
- Deploy model for prediction
- Monitor model performance

## 6.2 2. Real-Time Sentiment Analysis on Social Media

**Objective:** Analyze and visualize the sentiment of social media posts (tweets) in real-time with robust data processing capabilities.

### 6.2.1 Detailed Steps and Technologies

#### 1. Data Collection System

a) Twitter API Integration:

- Set up Twitter Developer account
- Implement OAuth authentication
- Create API rate limiting handling
- Set up error handling and retry logic
- Implement streaming connection management

b) Data Collection Pipeline:

- Create tweet filtering by keywords
- Implement language detection
- Set up geolocation tracking
- Create user metadata collection
- Implement hashtag tracking

c) Real-time Data Validation:

- Create input data validation rules
- Implement duplicate detection
- Set up data quality checks
- Create error logging system
- Implement data correction procedures

#### 2. Data Streaming Architecture

a) Apache Kafka Setup:

- Configure Kafka clusters
- Create topic partitioning strategy

- Set up consumer groups
- Implement message serialization
- Configure retention policies

b) Stream Processing:

- Implement Kafka Streams for processing
- Create real-time filtering logic
- Set up stream aggregations
- Implement windowing operations
- Create fault tolerance handling

### **3. Data Cleaning Pipeline**

a) Text Preprocessing:

- Remove URLs and special characters
- Clean hashtags and mentions
- Handle emoji conversions
- Implement language-specific cleaning
- Create text normalization procedures

b) Data Standardization:

- Implement timestamp standardization
- Create user ID normalization
- Set up location data standardization
- Implement metadata formatting
- Create data validation checks

### **4. Natural Language Processing**

a) Text Processing:

- Implement tokenization
- Create stop word removal
- Set up lemmatization/stemming

- Implement part-of-speech tagging
  - Create named entity recognition
- b) Sentiment Analysis:
- Implement VADER sentiment analyzer
  - Create TextBlob integration
  - Set up custom sentiment rules
  - Implement sentiment score normalization
  - Create confidence scoring

## **5. Real-time Processing System**

- a) Stream Processing Pipeline:
- Create real-time processing workers
  - Implement parallel processing
  - Set up batch processing for backlog
  - Create processing queue management
  - Implement error handling

- b) Performance Optimization:

- Implement caching strategies
- Create load balancing
- Set up resource scaling
- Implement performance monitoring
- Create optimization feedback loops

## **6. Data Storage and Management**

- a) Raw Data Storage:
- Implement tweet storage procedures
  - Create metadata management
  - Set up historical data archiving
  - Implement data compression

- Create storage optimization

b) Processed Data Management:

- Create sentiment results storage
- Implement aggregation storage
- Set up trend data management
- Create data access layers
- Implement data lifecycle management

## 7. Analysis and Aggregation

a) Real-time Analytics:

- Create rolling sentiment averages
- Implement trend detection
- Set up anomaly detection
- Create volume analysis
- Implement topic clustering

b) Statistical Analysis:

- Create sentiment distribution analysis
- Implement correlation studies
- Set up significance testing
- Create predictive modeling
- Implement pattern recognition

## 8. Visualization System

a) Dashboard Development (Plotly Dash):

- Create real-time sentiment graphs
- Implement trend visualizations
- Set up geographic visualizations
- Create interactive filters
- Implement custom visualizations

b) Reporting System:

- Create automated report generation
- Implement custom metrics calculation
- Set up visualization templates
- Create export functionality

### **6.3 3. Sales Forecasting for an E-commerce Platform**

**Objective:** Develop a comprehensive sales forecasting system to predict future sales patterns using historical data with robust data science and real-time updating capabilities.

#### **6.3.1 Detailed Steps and Technologies**

##### **1. Data Storage Architecture**

a) Data Organization:

- Design file structure for products catalog management
- Organize sales transactions data with time-series optimization
- Structure customer data for buyer behavior analysis
- Set up inventory management data files
- Create forecast results storage
- Implement seasonal patterns data files

b) Data Performance Optimization:

- Configure efficient data indexing strategies
- Implement data partitioning by date ranges
- Optimize data access patterns
- Create efficient data structures
- Implement data caching mechanisms

c) Data Management:

- Set up automated backup procedures
- Create data archiving policies
- Implement monitoring and alerting

- Configure data redundancy
- Create disaster recovery procedures
- Set up data retention policies

## **2. Data Collection and Integration**

a) Historical Data Collection:

- Extract historical sales data
- Collect product catalog information
- Gather customer purchase histories
- Import inventory movement data
- Collect pricing history
- Extract promotional campaign data

b) Real-time Data Collection:

- Set up real-time sales tracking
- Implement inventory level monitoring
- Create price change tracking
- Monitor customer behavior
- Track website analytics
- Implement promotional activity tracking

c) Data Integration:

- Create data flow automation
- Set up source system connections
- Implement data transformation rules
- Create error handling procedures
- Set up data validation checks
- Implement logging and monitoring

## **3. ETL Pipeline Development**

a) Data Extraction:

- Create data connectors

- Implement API integrations
- Set up file system readers
- Create data validation rules
- Implement error handling
- Set up extraction scheduling

b) Data Transformation:

- Implement data cleaning rules
- Create data standardization procedures
- Set up data enrichment processes
- Create calculation pipelines
- Implement data quality checks
- Create transformation logging

c) Data Loading:

- Create loading procedures
- Implement transaction management
- Set up error handling
- Create loading validation
- Implement performance optimization
- Set up loading monitoring

#### **4. Data Preprocessing and Cleaning**

a) Data Cleaning:

- Handle missing values
- Remove duplicates
- Clean outliers
- Standardize formats
- Validate data types
- Implement data quality rules

b) Feature Engineering:

- Create time-based features
- Generate seasonal indicators
- Calculate moving averages
- Create lag features
- Implement interaction features
- Generate categorical encodings

## **5. Time Series Analysis**

a) Pattern Analysis:

- Identify seasonal patterns
- Analyze trends
- Detect cyclical components
- Study irregular variations
- Analyze special events impact
- Create pattern documentation

b) Statistical Analysis:

- Perform decomposition analysis
- Calculate correlation studies
- Implement statistical tests
- Create distribution analysis
- Generate summary statistics
- Perform variance analysis

## **6. Forecasting Model Development**

a) Model Selection and Implementation:

- Implement ARIMA models
- Create Prophet forecasting
- Set up exponential smoothing
- Implement machine learning models
- Create ensemble methods

- Set up model comparison framework

b) Model Training:

- Create training pipelines
- Implement cross-validation
- Set up hyperparameter tuning
- Create model validation
- Implement performance testing
- Set up model versioning

## 7. Model Evaluation System

a) Performance Metrics:

- Calculate MAE (Mean Absolute Error)
- Implement RMSE calculations
- Create MAPE analysis
- Set up accuracy metrics
- Implement bias checking
- Create confidence intervals

b) Model Validation:

- Create validation procedures
- Implement backtesting
- Set up cross-validation
- Create benchmark comparisons
- Implement scenario testing
- Set up sensitivity analysis

## 8. Real-time Prediction System

a) Prediction Pipeline:

- Create real-time forecasting
- Implement model serving

- Set up prediction scheduling
- Create update procedures
- Implement error handling
- Set up monitoring system

b) Performance Optimization:

- Implement caching
- Create load balancing
- Set up resource scaling
- Implement request queuing
- Create performance monitoring
- Set up optimization procedures

## **9. Visualization and Reporting**

a) Dashboard Development (Flask/Django):

- Create forecast visualizations
- Implement trend displays
- Set up comparative analysis views
- Create performance metrics display
- Implement interactive features
- Set up customizable reports

b) Reporting System:

- Create automated reports
- Implement export functionality
- Set up scheduling system
- Create custom report templates
- Implement notification system
- Set up report distribution